

anomaly_detection_using_clustering_algorithms

April 20, 2026

1 Anomaly Detection using Clustering Algorithms

This notebook performs anomaly detection on network-intrusion data using three unsupervised clustering algorithms:

- KMeans
- Agglomerative Clustering
- DBSCAN

The data is taken from the existing `AnomalyDetection/data` folder. The `class` label is **not used for clustering**; it is used only after clustering to evaluate how well each algorithm separated normal and anomalous traffic.

1.1 1. Imports and Settings

```
[1]: import os
from pathlib import Path
import warnings

os.environ.setdefault("MPLCONFIGDIR", "/tmp")
warnings.filterwarnings("ignore", category=FutureWarning)

import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import seaborn as sns
from IPython.display import display
from scipy.cluster.hierarchy import dendrogram, linkage

from sklearn.cluster import AgglomerativeClustering, DBSCAN, KMeans
from sklearn.decomposition import PCA
from sklearn.metrics import (
    calinski_harabasz_score,
    ConfusionMatrixDisplay,
    accuracy_score,
    classification_report,
    confusion_matrix,
    davies_bouldin_score,
    f1_score,
```

```

precision_recall_fscore_support,
silhouette_score,
)
from sklearn.neighbors import NearestNeighbors
from sklearn.preprocessing import StandardScaler

RANDOM_STATE = 42
SAMPLE_PER_CLASS = 1500
BW_HATCHES = ["///", "\\\\", "xx", "..", "++", "oo"]
BW_COLORS = ["white", "#E6E6E6", "#BFBFBF", "#8C8C8C", "#666666", "#D9D9D9"]
BW_MARKERS = ["o", "s", "^", "D", "x", "P", "*"]
BW_LINESTYLES = ["-", "--", "-.", ":"]

sns.set_theme(style="whitegrid")
plt.rcParams.update({
    "figure.dpi": 120,
    "axes.edgecolor": "black",
    "axes.labelcolor": "black",
    "xtick.color": "black",
    "ytick.color": "black",
    "text.color": "black",
})
pd.set_option("display.max_columns", 120)
pd.set_option("display.max_colwidth", 120)

```

1.2 1.1 Helper Functions

```

[2]: def add_bar_labels(ax, horizontal=False):
    for patch in ax.patches:
        value = patch.get_width() if horizontal else patch.get_height()
        if np.isnan(value):
            continue
        if horizontal:
            ax.annotate(
                f"{value:,.0f}",
                (value, patch.get_y() + patch.get_height() / 2),
                ha="left",
                va="center",
                xytext=(4, 0),
                textcoords="offset points",
                fontsize=9,
            )
        else:
            ax.annotate(
                f"{value:,.0f}",
                (patch.get_x() + patch.get_width() / 2, value),
                ha="center",
            )

```

```

        va="bottom",
        xytext=(0, 3),
        textcoords="offset points",
        fontsize=9,
    )

def plot_count_bar(counts, title, xlabel="", ylabel="", horizontal=False,
    ↪top_n=None):
    plot_counts = counts.copy()
    if top_n is not None:
        plot_counts = plot_counts.head(top_n)
    fig, ax = plt.subplots(figsize=(9, max(4, 0.35 * len(plot_counts)))) if
    ↪horizontal else (8, 4.5))
    colors = [BW_COLORS[i % len(BW_COLORS)] for i in range(len(plot_counts))]
    if horizontal:
        bars = ax.barh(plot_counts.index.astype(str), plot_counts.values,
    ↪color=colors, edgecolor="black", linewidth=1.0)
        ax.invert_yaxis()
        ax.set_xlabel(xlabel)
        ax.set_ylabel(ylabel)
        ax.margins(x=0.15)
    else:
        bars = ax.bar(plot_counts.index.astype(str), plot_counts.values,
    ↪color=colors, edgecolor="black", linewidth=1.0)
        ax.set_xlabel(xlabel)
        ax.set_ylabel(ylabel)
        ax.tick_params(axis="x", rotation=25)
        ax.margins(y=0.15)
    for idx, bar in enumerate(bars):
        bar.set_hatch(BW_HATCHES[idx % len(BW_HATCHES)])
    ax.set_title(title)
    add_bar_labels(ax, horizontal=horizontal)
    plt.tight_layout()
    plt.show()

def get_data_root():
    candidates = [
        Path("AnomalyDetection/data"),
        Path("../data"),
        Path.cwd() / "AnomalyDetection" / "data",
    ]
    return next(path for path in candidates if path.exists())

def stratified_sample(data, label_col, n_per_class=SAMPLE_PER_CLASS):

```

```

    parts = []
    for label_value, group in data.groupby(label_col):
        parts.append(group.sample(min(len(group), n_per_class),
        ↪random_state=RANDOM_STATE))
    return pd.concat(parts).sample(frac=1, random_state=RANDOM_STATE).
    ↪reset_index(drop=True)

def preprocess_features(data, label_col):
    feature_data = data.drop(columns=[label_col]).copy()
    categorical_cols = feature_data.select_dtypes(include=["object", "string",
    ↪"category"]).columns.tolist()
    processed = pd.get_dummies(feature_data, columns=categorical_cols,
    ↪drop_first=False)
    processed = processed.apply(pd.to_numeric, errors="coerce").replace([np.
    ↪inf, -np.inf], np.nan).fillna(0)
    scaler = StandardScaler()
    scaled = scaler.fit_transform(processed)
    return processed, scaled, scaler

def map_clusters_to_anomaly(cluster_labels, y_true):
    mapped = np.zeros_like(y_true, dtype=int)
    mapping_rows = []
    for cluster_id in sorted(pd.Series(cluster_labels).unique()):
        mask = cluster_labels == cluster_id
        if cluster_id == -1:
            mapped_label = 1
            rule = "DBSCAN noise -> anomaly"
        else:
            counts = pd.Series(y_true[mask]).value_counts()
            mapped_label = int(counts.idxmax())
            rule = "majority label for evaluation"
        mapped[mask] = mapped_label
        mapping_rows.append(
            {
                "cluster": cluster_id,
                "mapped_prediction": "anomaly" if mapped_label == 1 else
    ↪"normal",
                "cluster_size": int(mask.sum()),
                "true_anomaly_rate": float(y_true[mask].mean()) if mask.sum()
    ↪else np.nan,
                "rule": rule,
            }
        )
    return mapped, pd.DataFrame(mapping_rows)

```

```

def internal_cluster_metrics(X, cluster_labels, ignore_noise=True):
    labels = np.asarray(cluster_labels)
    if ignore_noise:
        valid_mask = labels != -1
        X_eval = X[valid_mask]
        labels_eval = labels[valid_mask]
    else:
        X_eval = X
        labels_eval = labels

    unique_labels = np.unique(labels_eval)
    if X_eval.shape[0] < 2 or len(unique_labels) < 2:
        return {
            "Silhouette Score": np.nan,
            "Davies-Bouldin Index": np.nan,
            "Calinski-Harabasz Index": np.nan,
            "Metric Notes": "Not enough non-noise clusters",
        }

    return {
        "Silhouette Score": silhouette_score(X_eval, labels_eval),
        "Davies-Bouldin Index": davies_bouldin_score(X_eval, labels_eval),
        "Calinski-Harabasz Index": calinski_harabasz_score(X_eval, labels_eval),
        "Metric Notes": "Noise ignored" if ignore_noise and np.any(labels == -1)
        ↪-1) else "All points used",
    }

def evaluate_clustering(name, cluster_labels, y_true, X_for_internal_metrics):
    y_pred, mapping_df = map_clusters_to_anomaly(cluster_labels, y_true)
    precision, recall, f1, _ = precision_recall_fscore_support(
        y_true,
        y_pred,
        average="weighted",
        zero_division=0,
    )
    result = {
        "Model": name,
        "Clusters": len(set(cluster_labels)) - (1 if -1 in cluster_labels else ↪
        ↪0),
        "Noise Points": int((cluster_labels == -1).sum()),
        "Accuracy": accuracy_score(y_true, y_pred),
        "Weighted Precision": precision,
        "Weighted Recall": recall,
        "Weighted F1": f1,
    }

```

```

        "Anomaly F1": f1_score(y_true, y_pred, zero_division=0),
    }
    result.update(internal_cluster_metrics(X_for_internal_metrics,
    ↪cluster_labels, ignore_noise=True))

    print(f"{name} cluster-to-label mapping")
    display(mapping_df)

    print(f"{name} classification report")
    display(pd.DataFrame(classification_report(y_true, y_pred,
    ↪target_names=["normal", "anomaly"], output_dict=True, zero_division=0)).T.
    ↪round(4))

    cm = confusion_matrix(y_true, y_pred, labels=[0, 1])
    disp = ConfusionMatrixDisplay(confusion_matrix=cm,
    ↪display_labels=["normal", "anomaly"])
    disp.plot(cmap="Greys", values_format="d", colorbar=False)
    plt.title(f"{name} Confusion Matrix")
    plt.tight_layout()
    plt.show()
    return result, y_pred

def plot_pca_scatter(pca_df, color_col, title):
    fig, ax = plt.subplots(figsize=(8, 5.5))
    labels = sorted(pca_df[color_col].unique(), key=lambda x: str(x))
    for idx, label_value in enumerate(labels):
        subset = pca_df[pca_df[color_col] == label_value]
        ax.scatter(
            subset["PC1"],
            subset["PC2"],
            s=18,
            facecolors="none" if idx % 2 == 0 else BW_COLORS[idx %
    ↪len(BW_COLORS)],
            edgecolors="black",
            marker=BW_MARKERS[idx % len(BW_MARKERS)],
            linewidths=0.7,
            alpha=0.75,
            label=str(label_value),
        )
    ax.set_title(title)
    ax.set_xlabel("PC1")
    ax.set_ylabel("PC2")
    ax.legend(title=color_col, bbox_to_anchor=(1.02, 1), loc="upper left")
    plt.tight_layout()
    plt.show()

```

```

def plot_metric_lines(df, x_col, y_cols, title, ylabel):
    fig, ax = plt.subplots(figsize=(8.5, 4.8))
    for idx, col in enumerate(y_cols):
        ax.plot(
            df[x_col],
            df[col],
            color=BW_COLORS[min(idx + 1, len(BW_COLORS) - 1)] if idx else_
↪ "black",
            marker=BW_MARKERS[idx % len(BW_MARKERS)],
            linestyle=BW_LINESTYLES[idx % len(BW_LINESTYLES)],
            linewidth=1.8,
            label=col,
        )
    ax.set_title(title)
    ax.set_xlabel(x_col)
    ax.set_ylabel(ylabel)
    ax.legend()
    plt.tight_layout()
    plt.show()

```

1.3 2. Load Dataset

```

[3]: DATA_ROOT = get_data_root()
raw_path = DATA_ROOT / "raw" / "Network_Intrusion_Detection" / "Train_data.csv"

df_raw = pd.read_csv(raw_path)
print("Using dataset:", raw_path.resolve())
print("Raw shape:", df_raw.shape)
display(df_raw.head())

```

Using dataset: /home/durgaumadev/security-for-data-science-lab/AnomalyDetection/data/raw/Network_Intrusion_Detection/Train_data.csv
Raw shape: (25192, 42)

	duration	protocol_type	service	flag	src_bytes	dst_bytes	land	\
0	0	tcp	ftp_data	SF	491	0	0	
1	0	udp	other	SF	146	0	0	
2	0	tcp	private	S0	0	0	0	
3	0	tcp	http	SF	232	8153	0	
4	0	tcp	http	SF	199	420	0	

	wrong_fragment	urgent	hot	num_failed_logins	logged_in	num_compromised	\
0	0	0	0	0	0	0	
1	0	0	0	0	0	0	
2	0	0	0	0	0	0	
3	0	0	0	0	1	0	
4	0	0	0	0	1	0	

	root_shell	su_attempted	num_root	num_file_creations	num_shells	\
0	0	0	0	0	0	
1	0	0	0	0	0	
2	0	0	0	0	0	
3	0	0	0	0	0	
4	0	0	0	0	0	

	num_access_files	num_outbound_cmds	is_host_login	is_guest_login	count	\
0	0	0	0	0	2	
1	0	0	0	0	13	
2	0	0	0	0	123	
3	0	0	0	0	5	
4	0	0	0	0	30	

	srv_count	serror_rate	srv_serror_rate	rerror_rate	srv_rerror_rate	\
0	2	0.0	0.0	0.0	0.0	
1	1	0.0	0.0	0.0	0.0	
2	6	1.0	1.0	0.0	0.0	
3	5	0.2	0.2	0.0	0.0	
4	32	0.0	0.0	0.0	0.0	

	same_srv_rate	diff_srv_rate	srv_diff_host_rate	dst_host_count	\
0	1.00	0.00	0.00	150	
1	0.08	0.15	0.00	255	
2	0.05	0.07	0.00	255	
3	1.00	0.00	0.00	30	
4	1.00	0.00	0.09	255	

	dst_host_srv_count	dst_host_same_srv_rate	dst_host_diff_srv_rate	\
0	25	0.17	0.03	
1	1	0.00	0.60	
2	26	0.10	0.05	
3	255	1.00	0.00	
4	255	1.00	0.00	

	dst_host_same_src_port_rate	dst_host_srv_diff_host_rate	\
0	0.17	0.00	
1	0.88	0.00	
2	0.00	0.00	
3	0.03	0.04	
4	0.00	0.00	

	dst_host_serror_rate	dst_host_srv_serror_rate	dst_host_rerror_rate	\
0	0.00	0.00	0.05	
1	0.00	0.00	0.00	
2	1.00	1.00	0.00	
3	0.03	0.01	0.00	

4	0.00	0.00	0.00
---	------	------	------

	dst_host_srv_error_rate	class
0	0.00	normal
1	0.00	normal
2	0.00	anomaly
3	0.01	normal
4	0.00	normal

1.4 3. Exploratory Data Analysis

```
[4]: print("Columns:", df_raw.columns.tolist())
print("\nMissing values")
display(pd.DataFrame({
    "missing": df_raw.isna().sum(),
    "missing_pct": (df_raw.isna().mean() * 100).round(2),
}).sort_values("missing", ascending=False).head(15))

class_counts = df_raw["class"].value_counts()
display(class_counts.rename_axis("class").reset_index(name="count"))
plot_count_bar(class_counts, title="Normal vs Anomaly Class Distribution",
               ylabel="Records")

for col in ["protocol_type", "flag", "service"]:
    print(f"Top values for {col}")
    display(df_raw[col].value_counts().head(12).rename_axis(col).
           reset_index(name="count"))
    plot_count_bar(df_raw[col].value_counts(), title=f"Top {col} Values",
                   xlabel="Records", horizontal=True, top_n=12)
```


```
Columns: ['duration', 'protocol_type', 'service', 'flag', 'src_bytes',
'dst_bytes', 'land', 'wrong_fragment', 'urgent', 'hot', 'num_failed_logins',
'logged_in', 'num_compromised', 'root_shell', 'su_attempted', 'num_root',
'num_file_creations', 'num_shells', 'num_access_files', 'num_outbound_cmds',
'is_host_login', 'is_guest_login', 'count', 'srv_count', 'error_rate',
'srv_error_rate', 'error_rate', 'srv_error_rate', 'same_srv_rate',
'diff_srv_rate', 'srv_diff_host_rate', 'dst_host_count', 'dst_host_srv_count',
'dst_host_same_srv_rate', 'dst_host_diff_srv_rate',
'dst_host_same_src_port_rate', 'dst_host_srv_diff_host_rate',
'dst_host_error_rate', 'dst_host_srv_error_rate', 'dst_host_error_rate',
'dst_host_srv_error_rate', 'class']
```

Missing values

	missing	missing_pct
duration	0	0.0
protocol_type	0	0.0
service	0	0.0


flag	0	0.0
src_bytes	0	0.0
dst_bytes	0	0.0
land	0	0.0
wrong_fragment	0	0.0
urgent	0	0.0
hot	0	0.0
num_failed_logins	0	0.0
logged_in	0	0.0
num_compromised	0	0.0
root_shell	0	0.0
su_attempted	0	0.0

	class	count
0	normal	13449
1	anomaly	11743




Top values for protocol_type

	protocol_type	count
0	tcp	20526
1	udp	3011
2	icmp	1655




Top values for flag

	flag	count
0	SF	14973
1	S0	7009
2	REJ	2216
3	RSTR	497
4	RSTO	304
5	S1	88
6	SH	43
7	RSTOS0	21
8	S2	21
9	S3	15
10	OTH	5



Top values for service

	service	count
0	http	8003
1	private	4351
2	domain_u	1820
3	smtp	1449
4	ftp_data	1396
5	eco_i	909
6	other	858
7	ecr_i	613
8	telnet	483
9	finger	366
10	ftp	345
11	auth	189



1.5 4. Sampling and Preprocessing

```
[5]: df = stratified_sample(df_raw, label_col="class", n_per_class=SAMPLE_PER_CLASS)
print("Sampled shape:", df.shape)
display(df["class"].value_counts().rename_axis("class").
        ↪reset_index(name="count"))

y_true = (df["class"] != "normal").astype(int).to_numpy()
X_processed, X_scaled, scaler = preprocess_features(df, label_col="class")

print("Processed feature shape:", X_processed.shape)
display(X_processed.head())
```

Sampled shape: (3000, 42)

	class	count
0	normal	1500
1	anomaly	1500

Processed feature shape: (3000, 109)

	duration	src_bytes	dst_bytes	land	wrong_fragment	urgent	hot	\
0	0	2444	329	0	0	0	0	
1	0	0	0	0	0	0	0	
2	0	10044	0	0	0	0	0	
3	0	1032	0	0	0	0	0	
4	0	163	1510	0	0	0	0	

	num_failed_logins	logged_in	num_compromised	root_shell	su_attempted	\
0	0	1	0	0	0	
1	0	0	0	0	0	
2	0	1	0	0	0	
3	0	0	0	0	0	
4	0	1	0	0	0	

	num_root	num_file_creations	num_shells	num_access_files	\
0	0	0	0	0	
1	0	0	0	0	
2	0	0	0	0	
3	0	0	0	0	
4	0	0	0	0	

	num_outbound_cmds	is_host_login	is_guest_login	count	srv_count	\
0	0	0	0	2	2	
1	0	0	0	225	25	
2	0	0	0	12	12	
3	0	0	0	511	511	
4	0	0	0	3	7	

	error_rate	srv_error_rate	rerror_rate	srv_rerror_rate	same_srv_rate	\
0	0.0	0.0	0.0	0.0	1.00	
1	1.0	1.0	0.0	0.0	0.11	
2	0.0	0.0	0.0	0.0	1.00	
3	0.0	0.0	0.0	0.0	1.00	
4	0.0	0.0	0.0	0.0	1.00	

	diff_srv_rate	srv_diff_host_rate	dst_host_count	dst_host_srv_count	\
0	0.00	0.00	255	157	
1	0.05	0.00	255	25	
2	0.00	0.00	18	18	
3	0.00	0.00	255	124	
4	0.00	0.29	3	255	

	dst_host_same_srv_rate	dst_host_diff_srv_rate	\
--	------------------------	------------------------	---

0	0.62	0.02
1	0.10	0.05
2	1.00	0.00
3	0.49	0.02
4	1.00	0.00

	dst_host_same_src_port_rate	dst_host_srv_diff_host_rate \
0	0.00	0.00
1	0.00	0.00
2	1.00	0.00
3	0.49	0.00
4	0.33	0.01

	dst_host_serror_rate	dst_host_srv_serror_rate	dst_host_rerror_rate \
0	0.00	0.0	0.0
1	1.00	1.0	0.0
2	0.00	0.0	0.0
3	0.02	0.0	0.0
4	0.00	0.0	0.0

	dst_host_srv_rerror_rate	protocol_type_icmp	protocol_type_tcp \
0	0.0	False	True
1	0.0	False	True
2	0.0	False	True
3	0.0	True	False
4	0.0	False	True

	protocol_type_udp	service_IRC	service_Z39_50	service_auth	service_bgp \
0	False	False	False	False	False
1	False	False	False	False	False
2	False	False	False	False	False
3	False	False	False	False	False
4	False	False	False	False	False

	service_courier	service_csnet_ns	service_ctf	service_daytime \
0	False	False	False	False
1	False	False	False	False
2	False	False	False	False
3	False	False	False	False
4	False	False	False	False

	service_discard	service_domain	service_domain_u	service_echo \
0	False	False	False	False
1	False	False	False	True
2	False	False	False	False
3	False	False	False	False
4	False	False	False	False

	service_eco_i	service_ecr_i	service_efs	service_exec	service_finger	\
0	False	False	False	False	False	
1	False	False	False	False	False	
2	False	False	False	False	False	
3	False	True	False	False	False	
4	False	False	False	False	False	

	service_ftp	service_ftp_data	service_gopher	service_hostnames	\
0	False	False	False	False	
1	False	False	False	False	
2	False	True	False	False	
3	False	False	False	False	
4	False	False	False	False	

	service_http	service_http_443	service_imap4	service_iso_tsap	\
0	False	False	False	False	
1	False	False	False	False	
2	False	False	False	False	
3	False	False	False	False	
4	True	False	False	False	

	service_klogin	service_kshell	service_ldap	service_link	service_login	\
0	False	False	False	False	False	
1	False	False	False	False	False	
2	False	False	False	False	False	
3	False	False	False	False	False	
4	False	False	False	False	False	

	service_mtp	service_name	service_netbios_dgm	service_netbios_ns	\
0	False	False	False	False	
1	False	False	False	False	
2	False	False	False	False	
3	False	False	False	False	
4	False	False	False	False	

	service_netbios_ssn	service_netstat	service_nnsp	service_nntp	\
0	False	False	False	False	
1	False	False	False	False	
2	False	False	False	False	
3	False	False	False	False	
4	False	False	False	False	

	service_ntp_u	service_other	service_pop_2	service_pop_3	\
0	False	False	False	False	
1	False	False	False	False	
2	False	False	False	False	
3	False	False	False	False	
4	False	False	False	False	

	service_printer	service_private	service_remote_job	service_rje	\
0	False	False	False	False	
1	False	False	False	False	
2	False	False	False	False	
3	False	False	False	False	
4	False	False	False	False	

	service_shell	service_smtp	service_sql_net	service_ssh	service_sunrpc	\
0	False	True	False	False	False	
1	False	False	False	False	False	
2	False	False	False	False	False	
3	False	False	False	False	False	
4	False	False	False	False	False	

	service_supdup	service_systat	service_telnet	service_time	\
0	False	False	False	False	
1	False	False	False	False	
2	False	False	False	False	
3	False	False	False	False	
4	False	False	False	False	

	service_urp_i	service_uucp	service_uucp_path	service_vmnet	\
0	False	False	False	False	
1	False	False	False	False	
2	False	False	False	False	
3	False	False	False	False	
4	False	False	False	False	

	service_whois	flag_REJ	flag_RSTO	flag_RSTR	flag_SO	flag_S1	flag_S2	\
0	False	False	False	False	False	False	False	
1	False	False	False	False	True	False	False	
2	False	False	False	False	False	False	False	
3	False	False	False	False	False	False	False	
4	False	False	False	False	False	False	False	


	flag_SF	flag_SH
0	True	False
1	False	False
2	True	False
3	True	False
4	True	False

1.6 5. PCA Visualization

```
[6]: pca = PCA(n_components=2, random_state=RANDOM_STATE)
X_pca = pca.fit_transform(X_scaled)
pca_df = pd.DataFrame(X_pca, columns=["PC1", "PC2"])
pca_df["true_label"] = np.where(y_true == 1, "anomaly", "normal")

print("Explained variance ratio:", np.round(pca.explained_variance_ratio_, 4))
plot_pca_scatter(pca_df, color_col="true_label", title="PCA View of True Normal_
↪vs Anomaly Labels")
```

Explained variance ratio: [0.0919 0.0591]



1.7 6. KMeans Clustering

```
[7]: k_values = range(2, 11)
kmeans_rows = []
for k in k_values:
    model = KMeans(n_clusters=k, n_init=20, random_state=RANDOM_STATE)
    labels = model.fit_predict(X_scaled)
    kmeans_rows.append({
        "k": k,
        "inertia": model.inertia_,
```

```

        "silhouette": silhouette_score(X_scaled, labels),
    })


kmeans_curve = pd.DataFrame(kmeans_rows)
display(kmeans_curve.round(4))
plot_metric_lines(kmeans_curve, "k", ["inertia"], "KMeans Elbow Curve",
    ↪ "Inertia")
plot_metric_lines(kmeans_curve, "k", ["silhouette"], "KMeans Silhouette by k",
    ↪ "Silhouette")


kmeans_model = KMeans(n_clusters=2, n_init=20, random_state=RANDOM_STATE)
kmeans_labels = kmeans_model.fit_predict(X_scaled)
pca_df["KMeans_cluster"] = kmeans_labels.astype(str)
plot_pca_scatter(pca_df, color_col="KMeans_cluster", title="KMeans Clusters on
    ↪ PCA View")

results = []
predictions = {}
kmeans_result, kmeans_pred = evaluate_clustering("KMeans", kmeans_labels,
    ↪ y_true, X_scaled)
results.append(kmeans_result)
predictions["KMeans"] = kmeans_pred

```

	k	inertia	silhouette
0	2	291026.8672	0.2203
1	3	272595.8227	0.2444
2	4	262986.8678	0.1651
3	5	259135.8047	0.1994
4	6	249211.9326	0.1817
5	7	241998.6411	0.2068
6	8	234677.1906	0.2162
7	9	228266.8111	0.2105
8	10	225509.4613	0.1957





KMeans cluster-to-label mapping


	cluster	mapped_prediction	cluster_size	true_anomaly_rate \
0	0	anomaly	1227	0.991850
1	1	normal	1773	0.159616

rule

0 majority label for evaluation
 1 majority label for evaluation

KMeans classification report

	precision	recall	f1-score	support
normal	0.8404	0.9933	0.9105	1500.0000
anomaly	0.9919	0.8113	0.8926	1500.0000
accuracy	0.9023	0.9023	0.9023	0.9023
macro avg	0.9161	0.9023	0.9015	3000.0000
weighted avg	0.9161	0.9023	0.9015	3000.0000



1.8 7. Agglomerative Clustering

```
[8]: dendrogram_sample_size = min(200, X_scaled.shape[0])
rng = np.random.default_rng(RANDOM_STATE)
dendrogram_idx = rng.choice(X_scaled.shape[0], size=dendrogram_sample_size,
                             ↪replace=False)
linked = linkage(X_scaled[dendrogram_idx], method="ward")

plt.figure(figsize=(12, 5))
dendrogram(linked, truncate_mode="level", p=4, color_threshold=0,
            ↪above_threshold_color="black")
plt.title("Hierarchical Clustering Dendrogram Sample")
plt.xlabel("Sample index")
plt.ylabel("Distance")
plt.tight_layout()
plt.show()



agg_model = AgglomerativeClustering(n_clusters=2, linkage="ward")
agg_labels = agg_model.fit_predict(X_scaled)
```

```

pca_df["Agglomerative_cluster"] = agg_labels.astype(str)
plot_pca_scatter(pca_df, color_col="Agglomerative_cluster",
    title="Agglomerative Clusters on PCA View")

agg_result, agg_pred = evaluate_clustering("Agglomerative", agg_labels, y_true,
    X_scaled)
results.append(agg_result)
predictions["Agglomerative"] = agg_pred

```



Agglomerative cluster-to-label mapping


	cluster	mapped_prediction	cluster_size	true_anomaly_rate	\
0	0	anomaly	1310	0.925191	
1	1	normal	1690	0.170414	

rule

- 0 majority label for evaluation
- 1 majority label for evaluation

Agglomerative classification report

	precision	recall	f1-score	support
normal	0.8296	0.9347	0.8790	1500.0000
anomaly	0.9252	0.8080	0.8626	1500.0000
accuracy	0.8713	0.8713	0.8713	0.8713
macro avg	0.8774	0.8713	0.8708	3000.0000
weighted avg	0.8774	0.8713	0.8708	3000.0000




1.9 8. DBSCAN Hyperparameter Exploration

```
[9]: min_samples = 10
neighbors = NearestNeighbors(n_neighbors=min_samples)
neighbors.fit(X_scaled)
distances, _ = neighbors.kneighbors(X_scaled)
k_distances = np.sort(distances[:, -1])

plt.figure(figsize=(9, 4.8))
plt.plot(np.arange(len(k_distances)), k_distances, color="black", linewidth=1.5)
plt.title(f"DBSCAN k-distance Plot (min_samples={min_samples})")
plt.xlabel("Sorted points")
plt.ylabel("Distance to kth nearest neighbor")
plt.tight_layout()
plt.show()

eps_candidates = np.quantile(k_distances, [0.75, 0.85, 0.90, 0.95, 0.98])
dbscan_rows = []
for eps in eps_candidates:
    labels = DBSCAN(eps=float(eps), min_samples=min_samples, n_jobs=-1).
    ↪fit_predict(X_scaled)
    pred, _ = map_clusters_to_anomaly(labels, y_true)
    dbscan_rows.append({
        "eps": float(eps),
        "clusters": len(set(labels)) - (1 if -1 in labels else 0),
        "noise_points": int((labels == -1).sum()),
        "noise_pct": (labels == -1).mean() * 100,
        "accuracy": accuracy_score(y_true, pred),
        "anomaly_f1": f1_score(y_true, pred, zero_division=0),
        **internal_cluster_metrics(X_scaled, labels, ignore_noise=True),
    })



dbscan_tuning = pd.DataFrame(dbscan_rows)
display(dbscan_tuning.round(4))
plot_metric_lines(dbscan_tuning, "eps", ["accuracy", "anomaly_f1"], "DBSCAN eps_
    ↪Tuning", "Score")
plot_metric_lines(dbscan_tuning, "eps", ["noise_pct"], "DBSCAN Noise Percentage_
    ↪by eps", "Noise %")
```

	eps	clusters	noise_points	noise_pct	accuracy	anomaly_f1	\
0	2.1874	48	623	20.7667	0.9123	0.9174	
1	4.7672	41	394	13.1333	0.7827	0.7431	
2	8.6115	31	223	7.4333	0.6910	0.5750	
3	11.5436	34	137	4.5667	0.6797	0.5413	
4	18.4213	1	51	1.7000	0.5090	0.0503	

	Silhouette Score	Davies-Bouldin Index	Calinski-Harabasz Index	\
0	0.6983	0.3491	1399.8455	
1	0.4420	0.5087	228.1475	
2	0.4077	0.6014	97.4711	
3	0.4765	0.5957	89.0824	
4	NaN	NaN	NaN	

	Metric Notes
0	Noise ignored
1	Noise ignored
2	Noise ignored
3	Noise ignored
4	Not enough non-noise clusters



1.10 9. DBSCAN Clustering

```
[ ]: best_eps = dbscan_tuning.sort_values("anomaly_f1", ascending=False).
      ↪iloc[0]["eps"]
      print("Selected eps from tuning table:", round(best_eps, 4))
```

```

dbscan_model = DBSCAN(eps=float(best_eps), min_samples=min_samples, n_jobs=-1)
dbscan_labels = dbscan_model.fit_predict(X_scaled)
pca_df["DBSCAN_cluster"] = dbscan_labels.astype(str)
plot_pca_scatter(pca_df, color_col="DBSCAN_cluster", title="DBSCAN Clusters and
↳Noise on PCA View")

dbscan_result, dbscan_pred = evaluate_clustering("DBSCAN", dbscan_labels,
↳y_true, X_scaled)
results.append(dbscan_result)
predictions["DBSCAN"] = dbscan_pred

```

1.11 10. Model Comparison

```

[11]: comparison_df = pd.DataFrame(results)
display(comparison_df.round(4))

print("Internal clustering metrics")
internal_metrics_df = comparison_df[
    [
        "Model",
        "Clusters",
        "Noise Points",
        "Silhouette Score",
        "Davies-Bouldin Index",
        "Calinski-Harabasz Index",
        "Metric Notes",
    ]
]
display(internal_metrics_df.round(4))

print("External evaluation metrics using known labels")
external_metrics_df = comparison_df[
    [
        "Model",
        "Accuracy",
        "Weighted Precision",
        "Weighted Recall",
        "Weighted F1",
        "Anomaly F1",
    ]
]
display(external_metrics_df.round(4))

plot_count_bar(
    comparison_df.set_index("Model")["Anomaly F1"],
    title="Clustering Algorithm Comparison by Anomaly F1",
    ylabel="Anomaly F1",

```

```

)

plot_count_bar(
    comparison_df.set_index("Model")["Weighted F1"],
    title="Clustering Algorithm Comparison by Weighted F1",
    ylabel="Weighted F1",
)

plot_count_bar(
    comparison_df.set_index("Model")["Silhouette Score"],
    title="Clustering Algorithm Comparison by Silhouette Score",
    ylabel="Silhouette Score",
)

plot_count_bar(
    comparison_df.set_index("Model")["Davies-Bouldin Index"],
    title="Clustering Algorithm Comparison by Davies-Bouldin Index",
    ylabel="Davies-Bouldin Index",
)

```

	Model	Clusters	Noise Points	Accuracy	Weighted Precision \
0	KMeans	2	0	0.9023	0.9161
1	Agglomerative	2	0	0.8713	0.8774
2	DBSCAN	48	623	0.9123	0.9186

	Weighted Recall	Weighted F1	Anomaly F1	Silhouette Score \
0	0.9023	0.9015	0.8926	0.2203
1	0.8713	0.8708	0.8626	0.2074
2	0.9123	0.9120	0.9174	0.6983

	Davies-Bouldin Index	Calinski-Harabasz Index	Metric Notes
0	2.6501	277.8625	All points used
1	2.7249	261.9802	All points used
2	0.3491	1399.8455	Noise ignored

Internal clustering metrics

	Model	Clusters	Noise Points	Silhouette Score \
0	KMeans	2	0	0.2203
1	Agglomerative	2	0	0.2074
2	DBSCAN	48	623	0.6983



	Davies-Bouldin Index	Calinski-Harabasz Index	Metric Notes
0	2.6501	277.8625	All points used
1	2.7249	261.9802	All points used
2	0.3491	1399.8455	Noise ignored



External evaluation metrics using known labels

Model	Accuracy	Weighted Precision	Weighted Recall	Weighted F1 \
-------	----------	--------------------	-----------------	---------------

0	KMeans	0.9023	0.9161	0.9023	0.9015
1	Agglomerative	0.8713	0.8774	0.8713	0.8708
2	DBSCAN	0.9123	0.9186	0.9123	0.9120

	Anomaly F1
0	0.8926
1	0.8626
2	0.9174





1.12 11. Final Prediction Table

```
[12]: prediction_table = pd.DataFrame({
    "actual": np.where(y_true == 1, "anomaly", "normal"),
    "KMeans": np.where(predictions["KMeans"] == 1, "anomaly", "normal"),
    "Agglomerative": np.where(predictions["Agglomerative"] == 1, "anomaly", "normal"),
    "DBSCAN": np.where(predictions["DBSCAN"] == 1, "anomaly", "normal"),
})
display(prediction_table.head(20))
```

	actual	KMeans	Agglomerative	DBSCAN
0	normal	normal	normal	normal
1	anomaly	anomaly	anomaly	anomaly
2	normal	normal	normal	normal
3	anomaly	normal	normal	anomaly
4	normal	normal	normal	normal
5	anomaly	anomaly	anomaly	anomaly
6	anomaly	anomaly	anomaly	anomaly
7	normal	normal	normal	normal
8	anomaly	anomaly	anomaly	anomaly
9	normal	normal	normal	normal
10	normal	normal	normal	normal
11	anomaly	anomaly	anomaly	anomaly
12	anomaly	normal	normal	normal
13	normal	normal	normal	normal
14	normal	normal	normal	normal
15	anomaly	anomaly	anomaly	anomaly
16	anomaly	anomaly	anomaly	anomaly
17	anomaly	anomaly	anomaly	anomaly
18	normal	normal	normal	normal
19	normal	normal	normal	normal